Reading Digits with Neural Networks

Todd Sierens[1,2]

[1]Perimeter Institute for Theoretical Physics
[2]University of Waterloo

February 16, 2016

Outline

Handwritten digits

Figure: 100 handwritten digits from the MNIST database provided as a training set by kaggle.com.

Neural Network

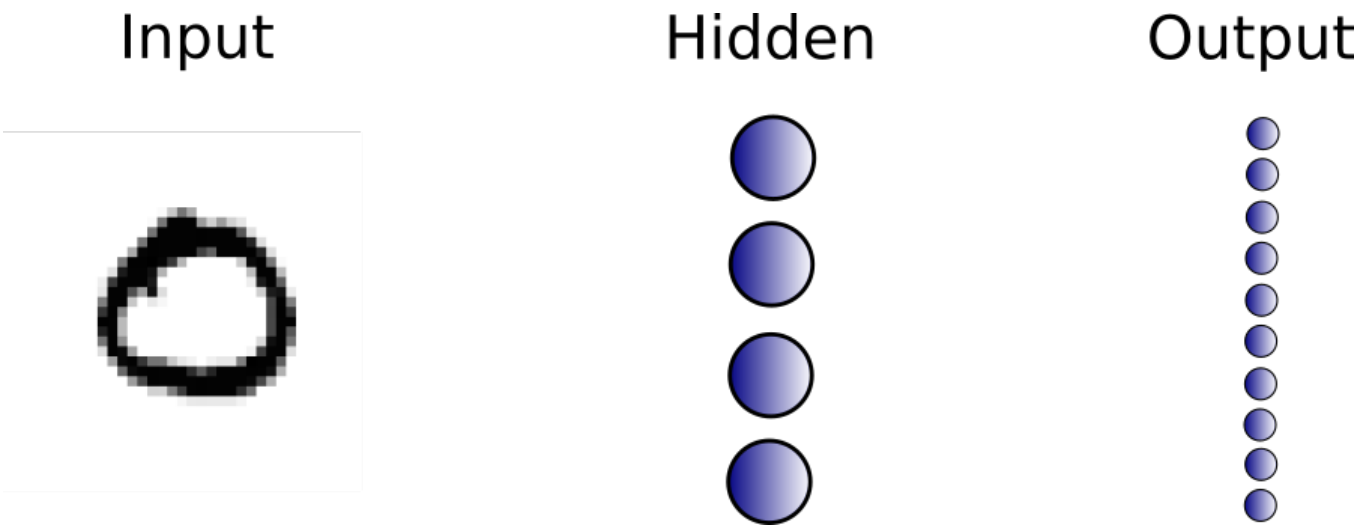

Neural Network

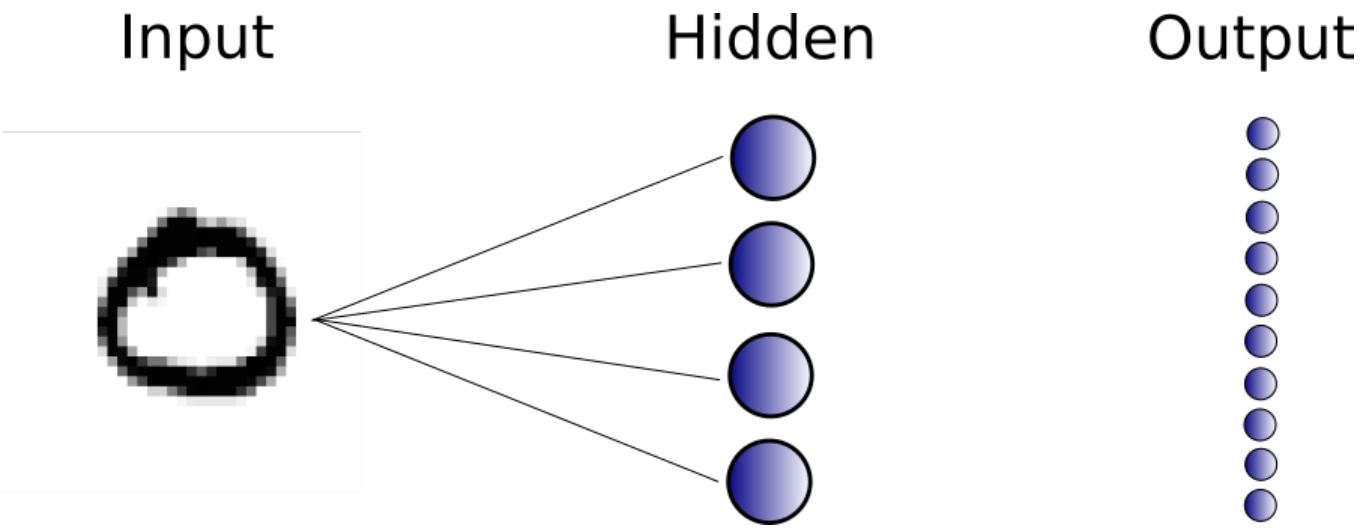

Neural Network

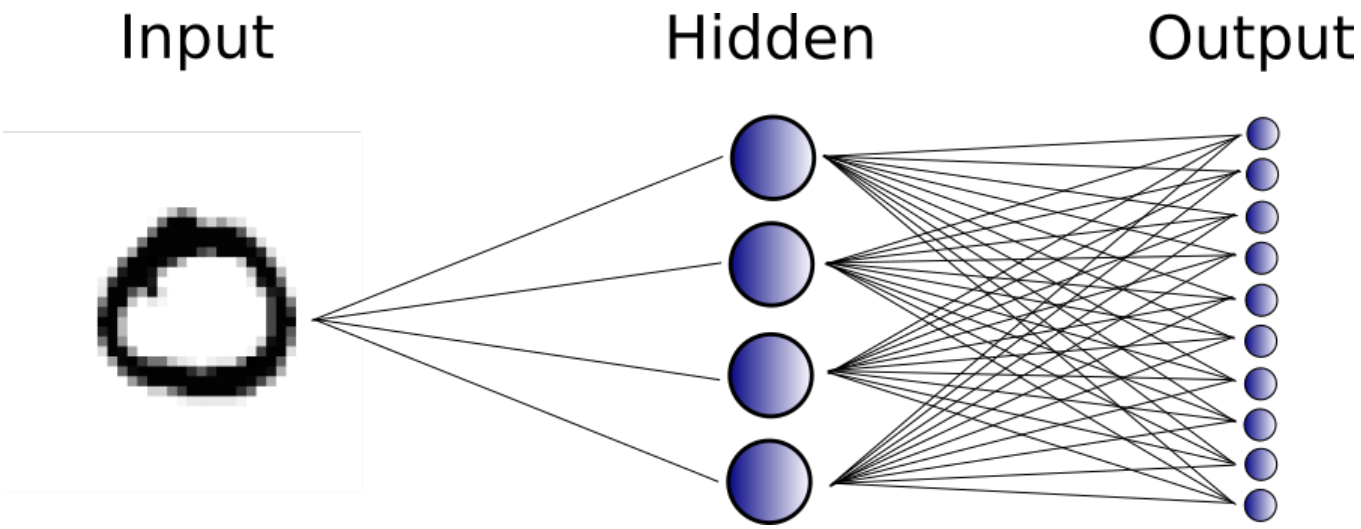

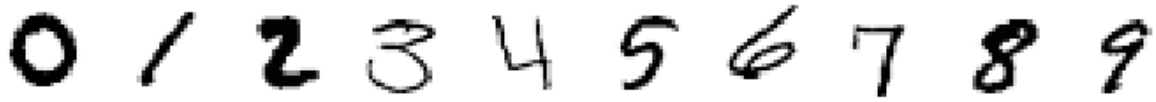

(a) Greyscale images of some of the training data.

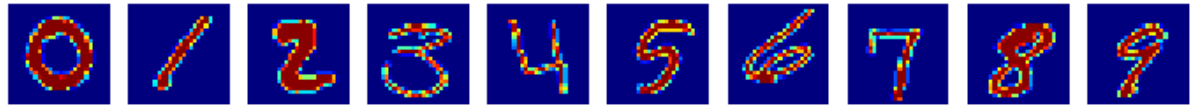

(b) Raw input layer of neural network.

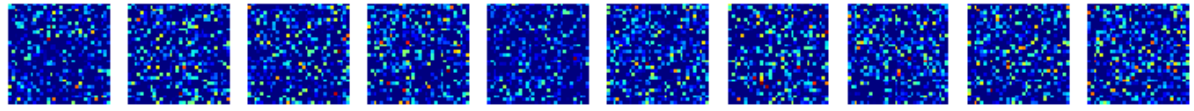

(c) Hidden layer of neural network.

Multilayer Perceptron

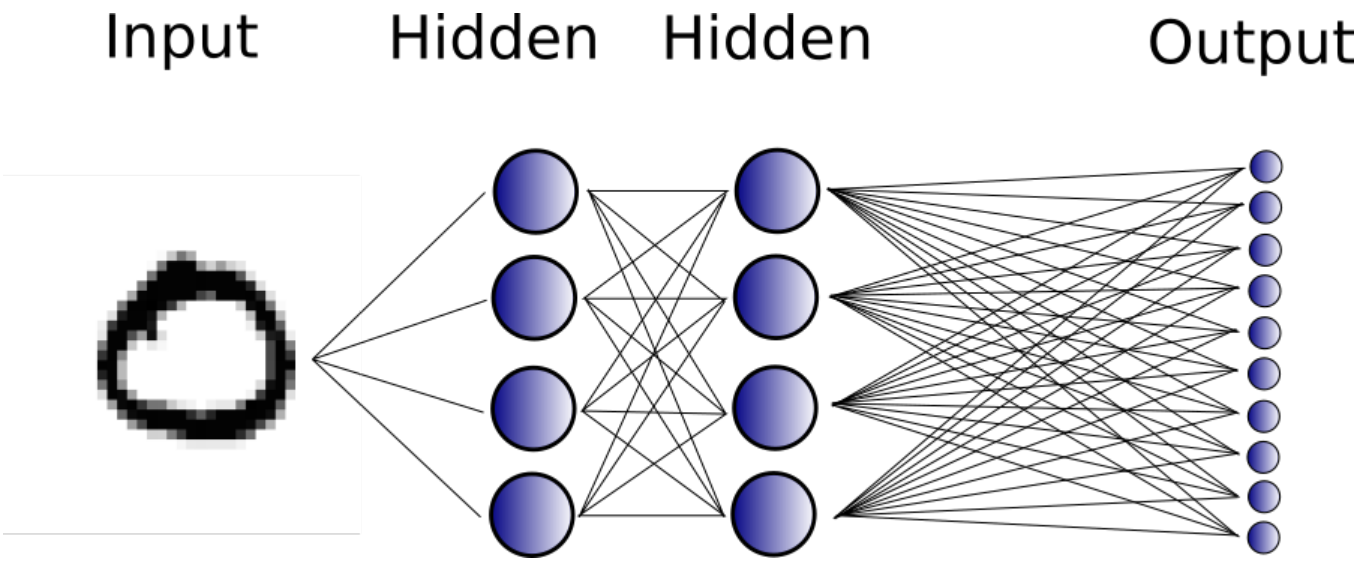

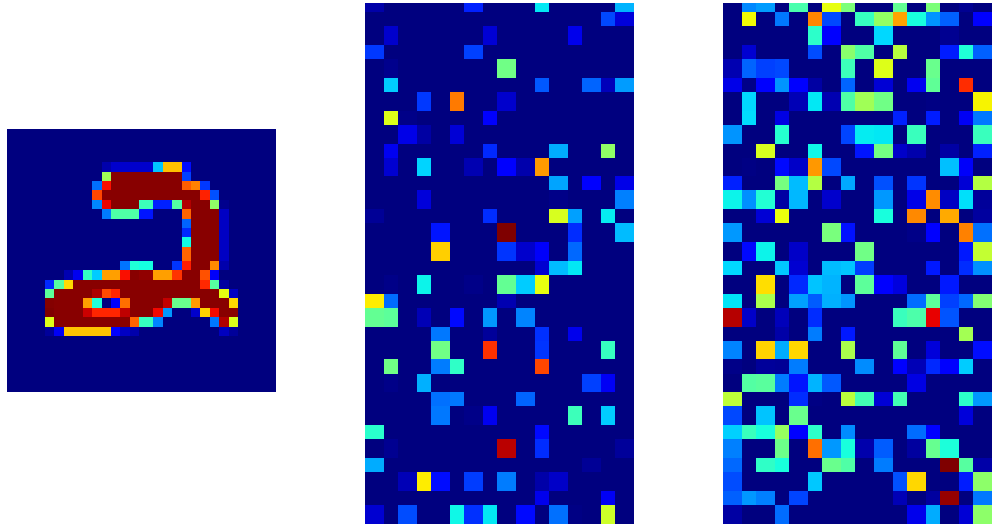

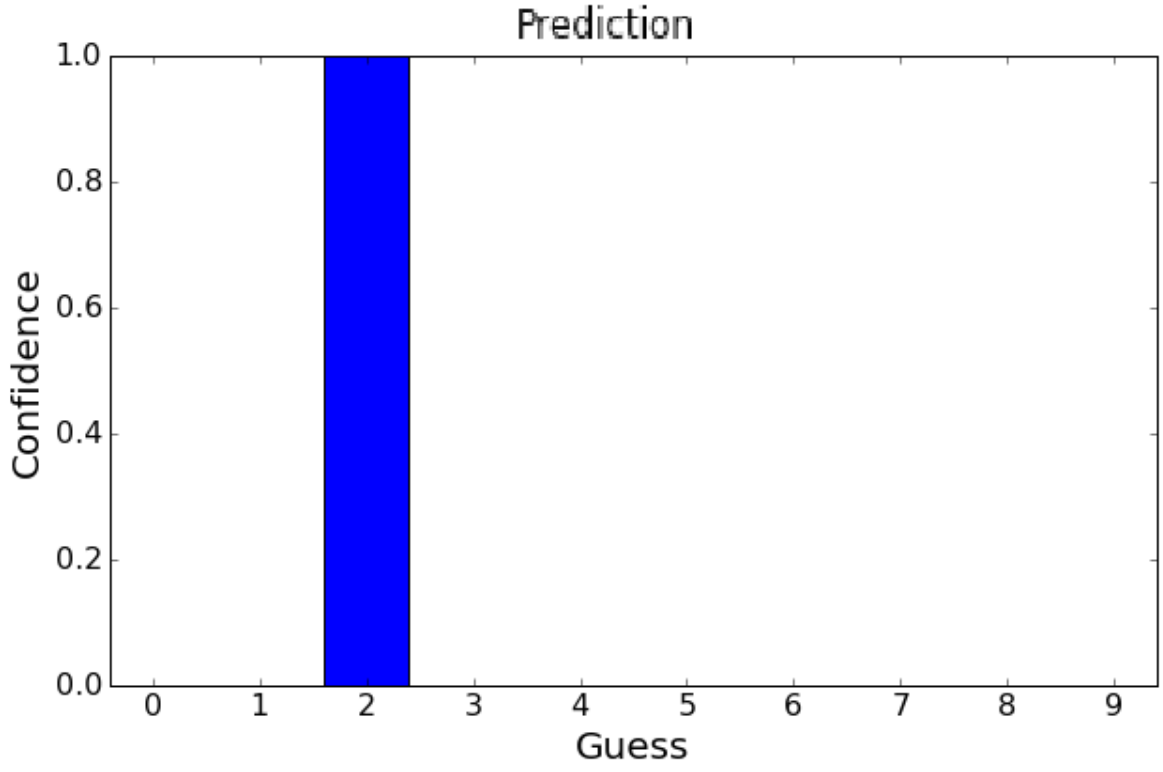
Prediction
1.0
0.8
0.6
0.4
0.2
0.0
Confidence
0
1
2
3
4
5
6
7
8
9
Guess

Convolutional Neural Network

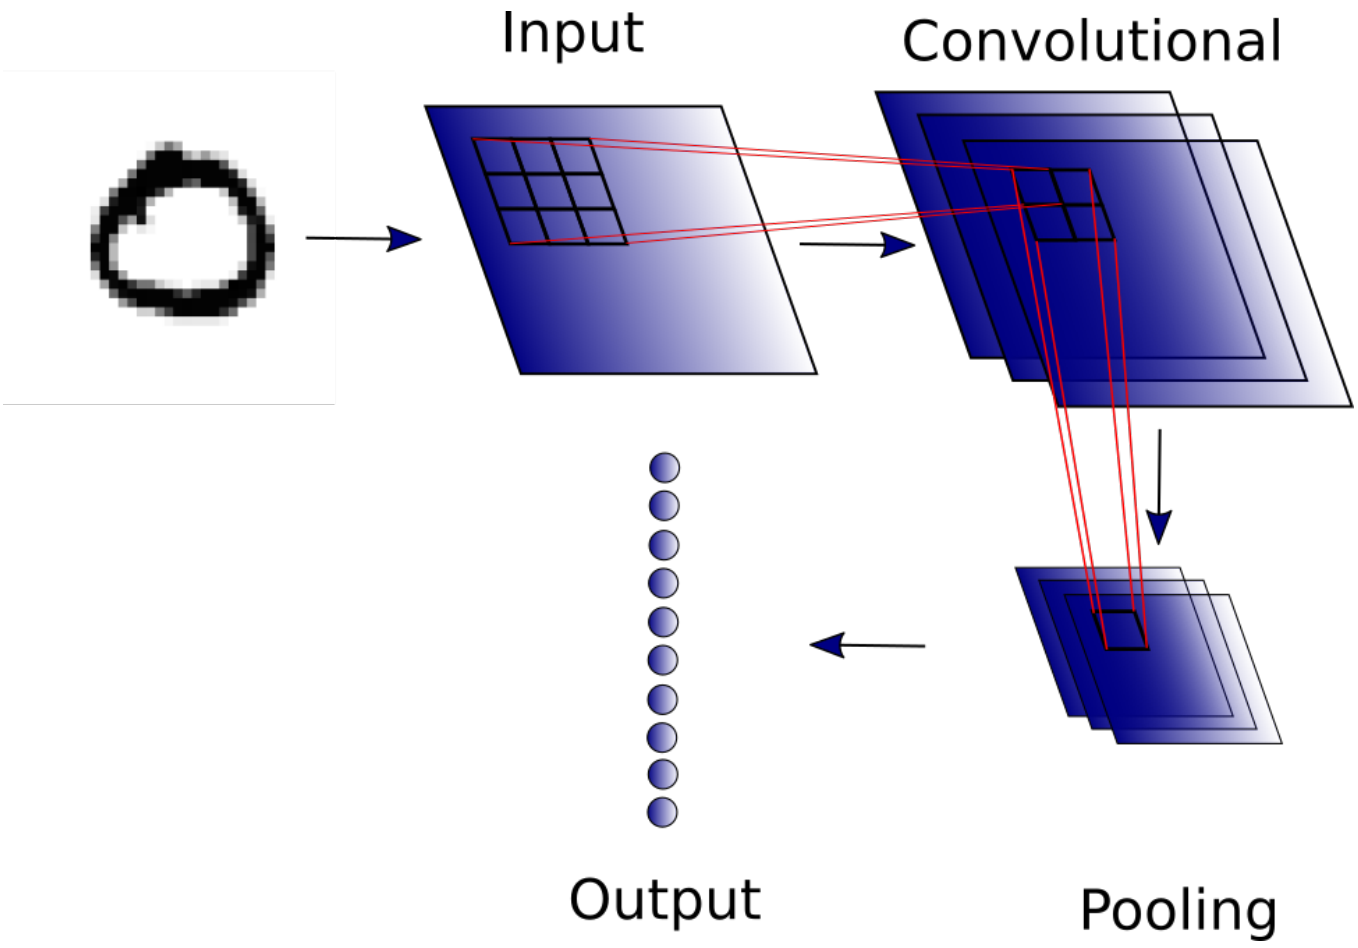

Input Layer

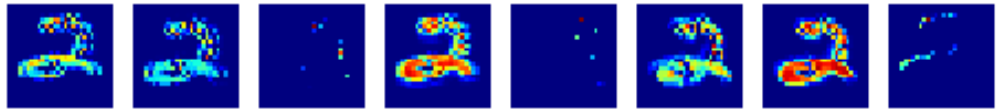

1st conv. layer

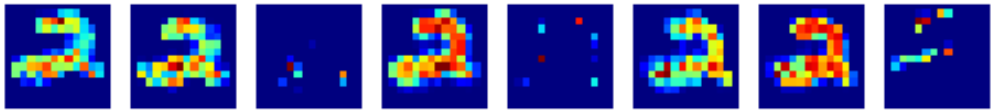

1st pooling layer

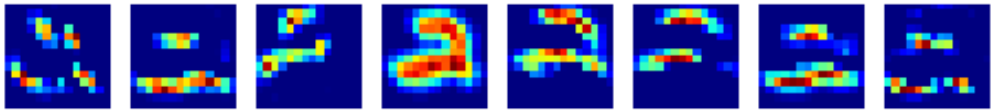

2nd conv. layer

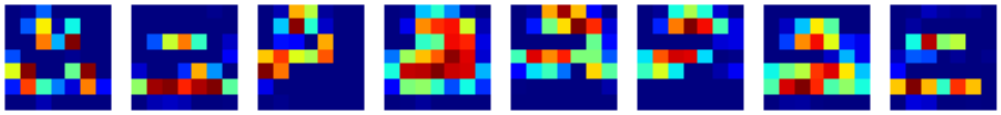

2nd pooling layer

Hidden layer

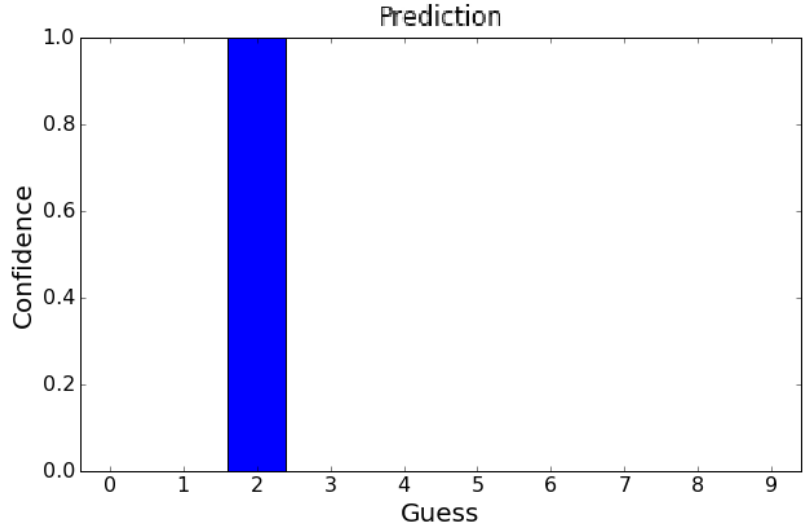

(a) prediction

- A deeper convolutional neural network (CNN)

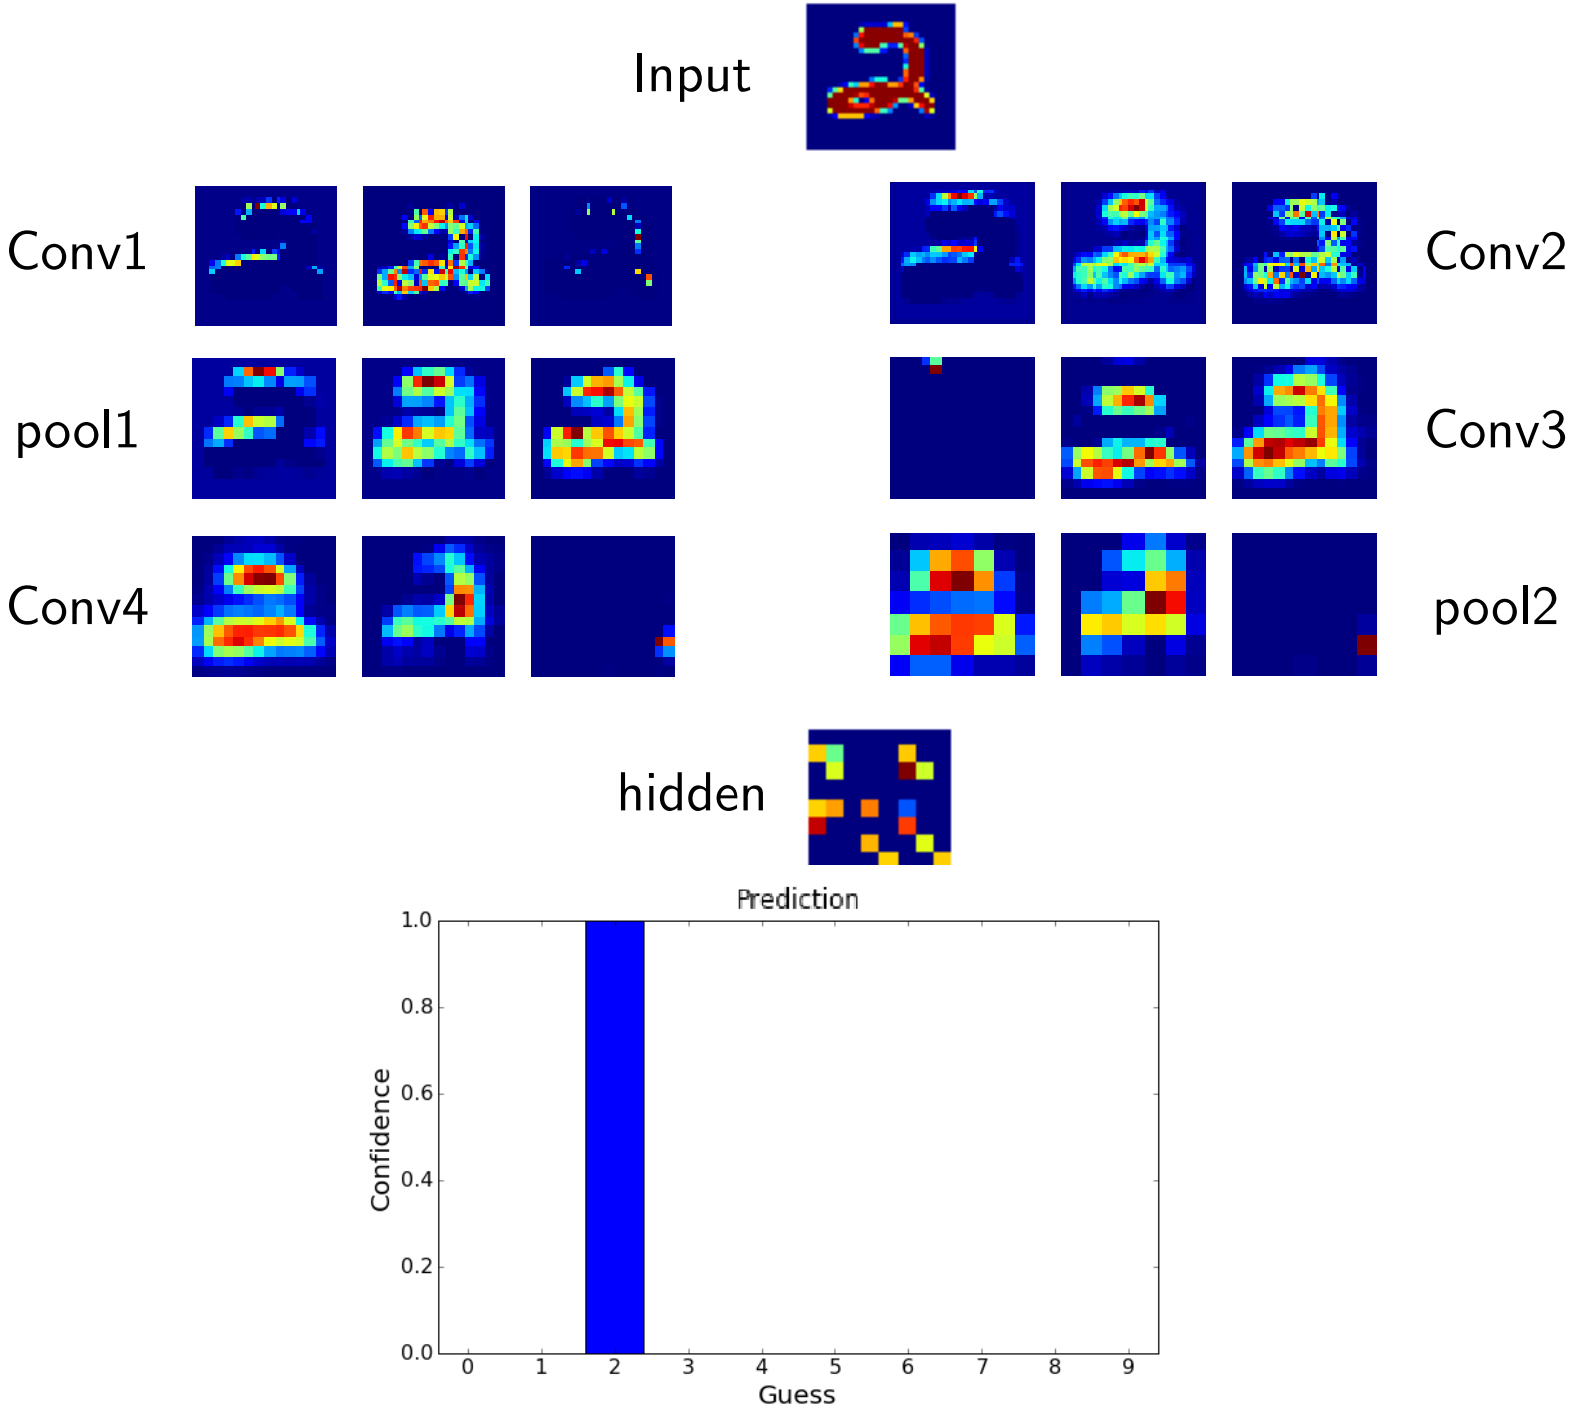
Input
Conv1
Conv2
pool1
Conv3
Conv4
pool2
hidden
Prediction
Confidence
1.0
0.8
0.6
0.4
0.2
0.0
0
1
2
3
4
5
6
7
8
9
Guess

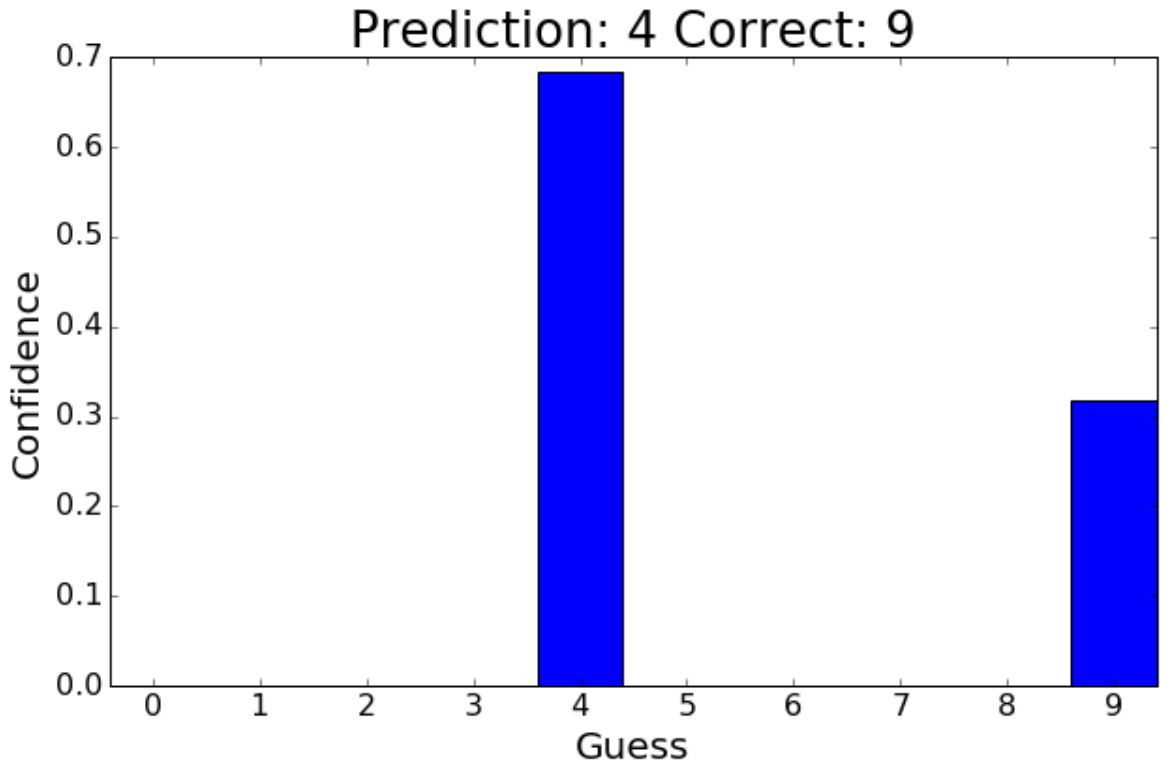
Prediction: 4 Correct: 9
Confidence
0.0
0.1
0.2
0.3
0.4
0.5
0.6
0.7
0
1
2
3
4
5
6
7
8
9
Guess

Performance Metrics

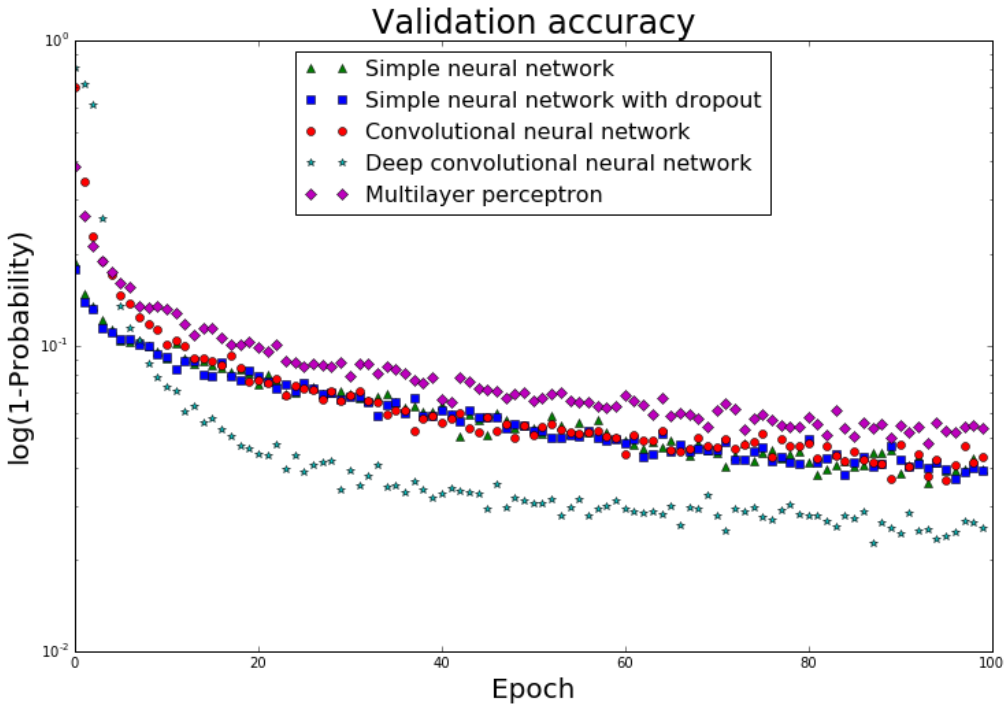

Figure: The validation accuracy of several models as a function of 'epoch'.

Performance Metrics

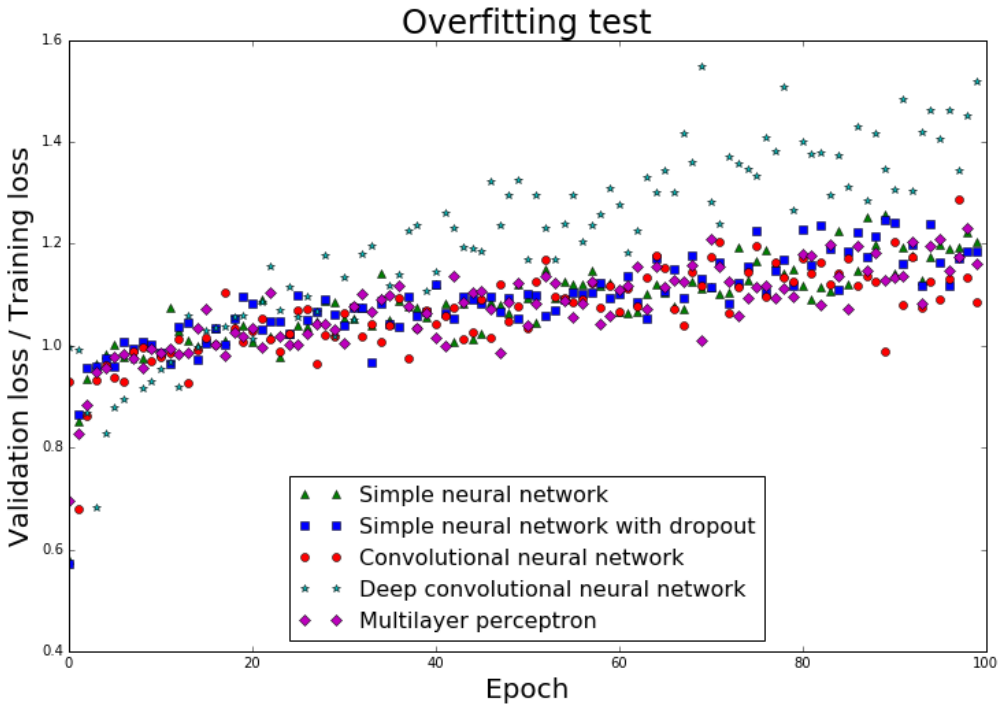

Figure: A measure of over-fitting. If this value is too far from 1 the model is over-fitting.

Performance Metrics

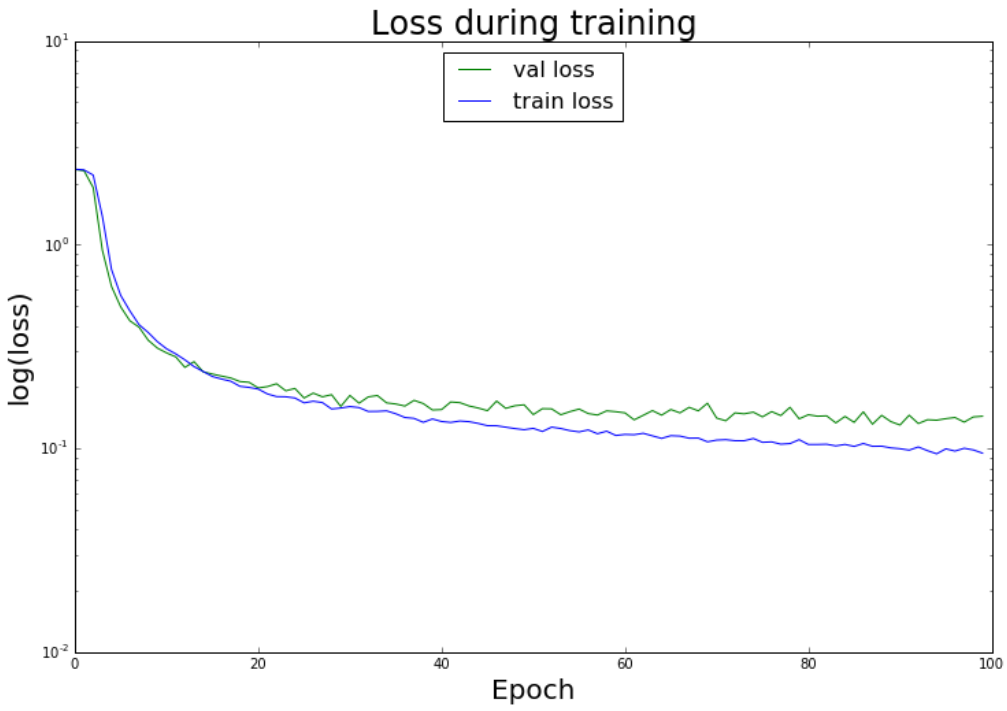

Figure: Validation loss and training loss for the deep CNN model.

Performance Metrics

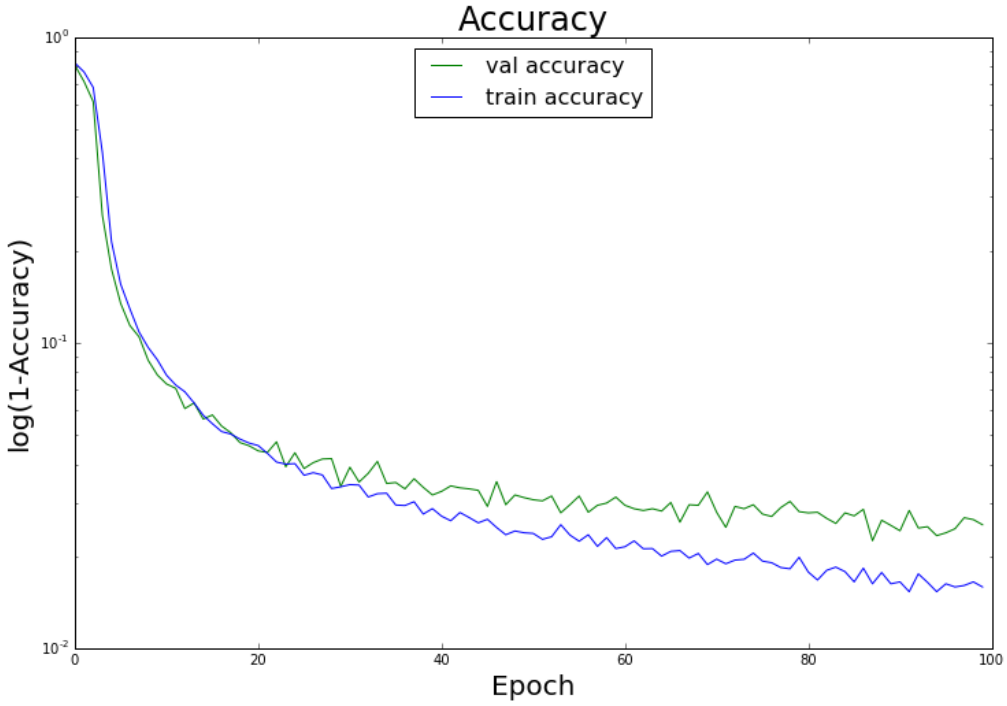

Figure: Validation accuracy and Training accuracy for the deep CNN model.

Performance Metrics

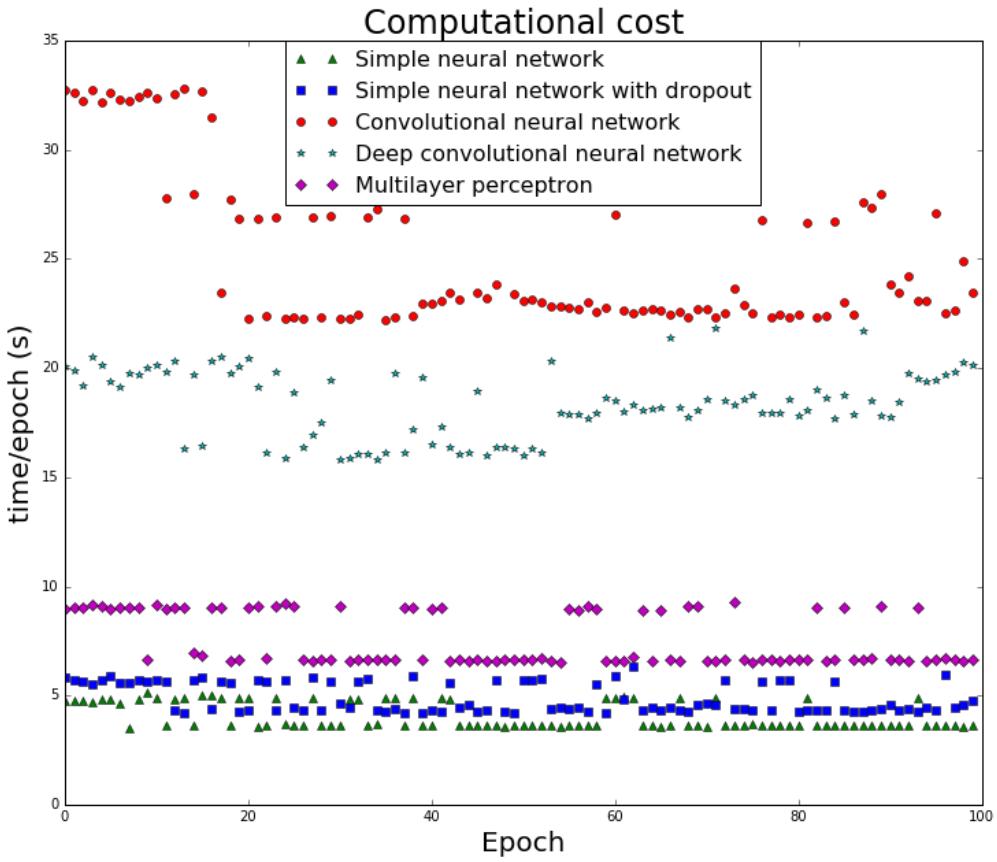

Figure: The amount of time each epoch required to execute. Since performance is dependent on many factors, only the relative time is meaningful.

The End

The End